

Localizing Generative Choices in Diffusion Models

a review

Panos Constantinides

School of Electrical and Computer Engineering, TU Crete, Greece

Source: <https://arxiv.org/abs/2604.06052v3>

Attention, May I Have Your Decision?

Localizing Generative Choices in Diffusion Models

Katarzyna Zaleska^{1*} Łukasz Popek¹ Monika Wysoczańska² Kamil Deja^{1,3}

¹Warsaw University of Technology ²valeo.ai ³IDEAS Research Institute

Problem statement [1-2]

Related work [2-3]

Methods [8-9]


Experimental results [2-3]

Limitations and future directions [2-3]

Implicit decisions and bias

Text-to-image diffusion produces high-quality results

Under-specified prompts require from the model to take biased decisions



Which parts of the model make these decisions?

How can we reduce the bias while preserving image quality?

Problem statement

Related work

Methods

Experimental results

Limitations and future directions

- **Latent space steering**

Techniques like *Latent Direction* and *Latent Editing* discover specific axes within the model's bottleneck and steer the generation during inference ([Li et al. 2024](#)).

- **Finetuning-based approaches**

These methods adapt the model's weights to minimize divergence from some prior distribution, effectively to favor more balanced outcomes.

- **Feature-based steering**

Sparse autoencoders identify and isolate the specific features responsible for biased attributes, allowing for more targeted steering ([Shi et al. 2025](#)).

- **Activation patching**

Establishes causal links by running the model on both clean and corrupted prompts, and swapping internal activations to determine which components are sufficient to alter the output ([Meng et al. 2022](#)).

Problem statement


Related work

Methods

Experimental results

Limitations and future directions

Implicit Choice-Modification (ICM)



Three stage process

1. Image generation and activation gathering
2. Pseudo-labeling via external classifier
3. Training layer-wise linear probes to quantify discriminability

Implicit decisions are made in self-attention layers.


There is no specific token to attend to, thus cross-attention layers perform worse in classification.

Explicit and implicit decisions


A classifier predicts *gender* based on activations of different attention layers of a UNet.



"A photo of a person"



"A photo of a man"



Decisions for **explicit** prompts are made in cross-attentions layers

Decisions for **implicit** prompts are made in self-attention layers

Now that we know which modules are responsible, how can we approach them?

Activation steering

1. Identify discriminating layers via ICM
2. Given the decision boundary $s_{l,t}$ for layer l and time-step t update activations as

$$H'_{l,t} = H_{l,t} + \alpha s_{l,t}$$

(at inference-time)

Fine-tuning

1. Construct a dataset of triplets $(x, p_{\text{gen}}, p_{\text{spec}})$
 x generated image from prompt p_{spec} but classified under p_{gen}
2. Fine-tune parameters of specific layers

(at training-time)

Problem statement

Related work

Methods

Experimental results

Limitations and future directions

Method	bias	image quality	similarity with reference image	alignment with prompt
	FD ↓	FID ↓	CLIP-I ↑	CLIP-T ↑
Original	0.564	120.06	-	0.6155
Finetuning [34]	0.050	161.47	0.8779	0.6095
ICM (Finetuning)	0.535	143.98	0.9187	0.6189
Latent Editing [16]	0.408	166.11	0.8253	0.6005
H-Distribution [26]	0.222	151.68	0.8475	0.6087
Latent Direction [19]	0.305	129.37	0.8058	0.6091
DIFFLENS [35]	0.046	112.83	0.8501	0.6090
ICM (Steering)	0.087	122.08	0.8500	0.6140


$$\text{FD} := \left\| \bar{p} - \mathbb{E}_{x \sim p_{\theta}(\mathbf{x})}(\mathbf{y}) \right\|_2$$

Trade-off

Bias reduction against image quality and prompt alignment

Significant bias reduction

Retains image quality




- Implicit decisions are indeed localized
- ICM avoids artifacts
- Average pooling successfully captures activations

Context and background remain the same

Layer selection matters

Scaling up to SDXL



Prompts describing a US president.

Positive class “Barack Obama”, all others are treated as negative.

- Debiasing via ICM scales well
- Preserves context
- Good balance across metrics

Layer Selection	Steering Success (%)	Avg. Confidence	CLIP-I $\uparrow$
Top 20 (Ours)	83.0%	0.917	0.893
Random 20	51.0%	0.838	0.932
All Layers	91.0%	0.879	0.779

Limitations

- The precision of the linear probes is tied to the accuracy and biases of the external classifier.
- Delicate tuning of the steering strength α .
- Capturing activations by average pooling may cause loss of detail.

Future directions

- Multi-factor control: pose, background etc.
- Faster steering for real-time generation.
- Investigate if self-attention is the universal module for implicit decisions in diffusion.